

UNIVERSITY OF SCIENCE AND TECHNOLOGY OF HANOI



BACHELOR THESIS

A Gamified E-Learning Platform for School Students

An Hoc - A Web-Based Gamified E-Learning Platform for Structured
Learning, Adaptive Practice, and Progress Tracking

Author:

DOAN Thuan An

BI11-002

Supervisor:

Dr. TRAN Giang Son

Hanoi - July 2026

DECLARATION

I hereby declare that this thesis titled *A Gamified E-Learning Platform for School Students* is prepared from the implemented Anhoc repository and the accompanying project documentation. All architectural descriptions, feature summaries, and implementation details presented in this report are derived from the inspected codebase, and external concepts or technologies are cited where appropriate.

To the best of my knowledge, this report does not intentionally reproduce uncited work as an original contribution. Any frameworks, libraries, documents, or external references used in the preparation of this thesis are acknowledged in the bibliography.

Hanoi, July 2026

DOAN Thuan An

ACKNOWLEDGEMENTS

This thesis was prepared from the current Anhoc web application repository. The document reflects the work embodied in the frontend, backend, database, and chatbot modules, as well as the supporting technical design notes maintained with the project.

I would like to acknowledge the contributors who shaped the platform's architecture, educational workflows, and deployment configuration, together with the maintainers of the open-source tools that made the system possible, including Next.js, React, Express, Prisma, PostgreSQL, and the supporting Python ecosystem used by the tutoring service.

(DOAN Thuan An)
Hanoi, July 2026

CONTENTS

Acknowledgements	3
List of Abbreviations	i
List of Tables	ii
List of Figures	iii
Abstract	v
I/ Introduction	1
1 Background	1
1.1 Statistical Context of E-Learning in Vietnam	1
2 Problem Statement	2
2.1 Advantages and Limitations of E-Learning Applications in Vietnam	2
3 Objectives	3
4 Scope	3
5 Contributions of the Project	4
6 Thesis Structure	4
II/ Objectives	5
1 Desired Features	5
2 Expected Outcome	6
III/ Materials and Methods	7
1 Requirements Analysis	7
1.1 Stakeholders and User Roles	7
1.2 Functional Requirements	7
1.3 Use Case Model	8
1.4 Detailed Use Case Diagrams	9
1.5 Use Case Descriptions	15
1.6 Non-Functional Requirements	19
1.7 Learning Content Requirements	20
1.8 Assessment Requirements	20

1.9	Administrative Requirements	20
1.10	Operational Assumptions	21
2	System Design	21
2.1	Architectural Overview	21
2.2	Deployment Design	22
2.3	Backend Design	23
2.4	Frontend Design	23
2.5	Security Model	23
2.6	Data Model	24
2.7	Entity-Relationship Diagrams	25
2.8	Question Snapshot Strategy	40
2.9	Progression Design	41
2.10	Sequence Diagrams for Main Use Cases	41
3	System Implementation	46
3.1	Frontend Implementation	46
3.2	Backend API Implementation	46
3.3	Question Generation and Grading	47
3.4	Attempt Lifecycle	47
3.5	Administrative Implementation	48
3.6	Configuration and Deployment	49
3.7	User Interface Presentation	50
IV/ Results and Discussion		51
1	Method of Evaluation	51
2	Representative User Scenarios	51
3	Observed Strengths	52
4	Current Limitations	52
5	Risks and Tradeoffs	52
6	Future Improvements	53
7	Discussion	53
V/ Conclusion & Perspective		54
Appendices		55

A.1	Build and Run Notes	55
A.2	Selected Environment Variables	55
A.3	Selected REST Endpoints	56
A.4	Application User Interface Screenshots	56
References		61

LIST OF ABBREVIATIONS

API	Application Programming Interface
BYOK	Bring Your Own Key
CORS	Cross-Origin Resource Sharing
JWT	JSON Web Token
LLM	Large Language Model
ORM	Object-Relational Mapping
RBAC	Role-Based Access Control
REST	Representational State Transfer
UI	User Interface
XP	Experience Points

LIST OF TABLES

3.1	Core functional requirements	8
3.10	Main persistent entities	24
3.27	Complete Database Schema — Table Summary (38 tables)	39
3.28	Main API groups	46
A.1	Representative environment variables	55
A.2	Example endpoints from the implemented API	56

LIST OF FIGURES

3.1	Use case diagram of the Anhoc platform	9
3.2	Use case diagram — Authentication and Account Management	10
3.3	Use case diagram — Learning System	11
3.4	Use case diagram — Practice and Assessment Engine	12
3.5	Use case diagram — AI Math Tutor Chatbot	13
3.6	Use case diagram — Gamification and Economy	14
3.7	Use case diagram — Subscription and Payments	15
3.8	Architecture design of the Anhoc platform	22
3.9	ERD Group 1 — User & Authentication tables	26
3.10	ERD Group 2 — Curriculum & Content tables	29
3.11	ERD Group 3 — Practice, Tests & Progress tables	32
3.12	ERD Group 4 — Gamification & Economy tables	35
3.13	ERD Group 5 — Subscriptions, Gaming and Notifications	37
3.14	Sequence diagram for registration or login	42
3.15	Sequence diagram for studying a lesson	42
3.16	Sequence diagram for completing lesson practice	43
3.17	Sequence diagram for taking a grade-level test	43
3.18	Sequence diagram for viewing progress and achievements	44
3.19	Sequence diagram for managing content and question templates	44
3.20	Sequence diagram for managing learning unit members and seats	45
3.21	Sequence diagram for managing roles, access control, and dashboard data	45
A.1	Student dashboard displaying experience points, active level, and unlocked achievements.	56
A.2	Grade-level learning pathway showing available lessons and bilingual subject content.	57
A.3	Interactive practice interface featuring a dynamically generated mathematics question and live evaluation feedback.	57
A.4	Achievements gallery displaying seeded badges and their associated experience point rewards.	58

A.5	Tutoring chatbot widget answering student queries with step-by-step mathematical reasoning.	58
A.6	Administrative overview panel showing system stats, active users, and lesson metrics.	59
A.7	Admin user management dashboard showing all registered users and their system roles.	59
A.8	Access control settings panel for configuring user seats and editing role-based permission profiles.	60

ABSTRACT

Anhoc is a web-based mathematics learning platform designed for structured school-level study. The implemented system combines bilingual lesson content, dynamically generated practice exercises, grade-level tests, progression tracking, administrative management, notifications, multiplayer and single-player math games, and an integrated tutoring chatbot. This thesis documents the design and implementation of Anhoc based on the current repository contents.

The software follows a service-oriented web architecture. A Next.js frontend provides the student and administrative interfaces. An Express backend exposes REST endpoints for authentication, lessons, tests, achievements, notifications, games, supervisor workflows, and subscription management. Prisma is used as the persistence layer, PostgreSQL stores the core application state, and a companion Python chatbot service extends the platform with conversational tutoring capabilities.

A notable implementation decision is the use of persisted question snapshots. Reusable templates define how questions are generated, while snapshots preserve the exact generated variables, accepted answers, and scoring outcomes for each attempt. This design supports reproducible grading, debugging, review, and reporting. The project also adopts dynamic role-based access control with subject-level restrictions, progression services for mastery and XP, and operational support for learn units and subscription-oriented platform management.

This thesis presents the problem context, requirements, system architecture, data model, implementation details, and an evaluation of the current solution. It also discusses limitations and practical improvements required to mature the platform further as an educational software system.

Keywords: mathematics learning platform, question generation, role-based access control, progress tracking, tutoring chatbot, web application architecture

I/ INTRODUCTION

1 Background

Digital learning systems are increasingly expected to do more than present static content. In mathematics education, students need repeated practice, immediate feedback, structured progression, and content that can be adapted to multiple grades and topics. A conventional content website can display lessons, but it does not by itself provide reusable question generation, mastery tracking, or administrative control over subject-specific learning material.

The Anhoc project addresses this gap by combining theory lessons, generated practice, formal test sessions, and gamified progression in a single web application. The project targets school-level mathematics learning and is organized around subjects, grades, and lessons. The implemented system allows a student to read theory content, start practice from lesson templates, submit answers, complete tests, and monitor progress through XP, levels, and achievements. At the same time, an administrator can manage lessons, question templates, user accounts, roles, and access permissions.

1.1 Statistical Context of E-Learning in Vietnam

The Vietnamese context makes digital learning both promising and uneven. DataReportal reported 78.44 million internet users in Vietnam in January 2024, equivalent to 79.1% of the population, together with 168.5 million active cellular mobile connections [10]. This creates a strong baseline for mobile-first learning applications. UNESCO also reported that about nine out of ten schools in Vietnam are connected to the internet and that national policy targets full school connectivity by 2025 [11]. EdTech Hub’s 2024 rapid scan further notes that the Vietnamese educational technology landscape already includes online learning platforms, learning apps, websites, and video-based resources, showing that e-learning is no longer a niche delivery mode in the country [14].

However, the same sources show that access and quality remain unequal. UNESCO highlights a very large home-access gap: 95% of children from the richest households benefit from internet access at home compared with only 18% of children from the poorest households. During the COVID-19 period, students from the poorest households were also 34% less likely to experience distance learning than students from the richest households, and there were roughly three students per computer on average [11]. A World Bank case note on Vietnam’s emergency distance-learning response adds that, in

less advantaged areas, only 10% to 20% of students and teachers had access to personal digital devices at home, even though some provinces still managed to reach about 80% of students through a combination of online and educational television lessons [13, 12].

2 Problem Statement

The main engineering problem solved by Anhoc is the design of a learning platform that satisfies four requirements at the same time:

- it must support reusable content across grades and subjects;
- it must generate exercises dynamically rather than storing every question instance manually;
- it must keep each attempt stable enough to be scored, reviewed, and audited later;
- it must provide operational control for administrators without hard-coding every permission into the user interface.

Many simple practice applications solve only one part of this problem. They either store fixed question banks, do not track long-term student progression, or rely on coarse user roles that become difficult to extend. Anhoc introduces a more structured approach by combining generated questions with persisted snapshots, dynamic role-based access control, and a progression subsystem.

2.1 Advantages and Limitations of E-Learning Applications in Vietnam

The statistical context above suggests several clear advantages for e-learning applications in Vietnam.

- **Large reachable user base:** high internet penetration and widespread mobile connectivity make app-based learning distribution technically feasible at national scale [10].
- **Institutional readiness:** the fact that most schools are already connected to the internet creates a practical foundation for blended learning, school-supported digital homework, and platform-assisted assessment [11].
- **Flexible continuity:** emergency distance-learning evidence shows that digital tools, especially when combined with broadcast channels, can preserve instructional continuity during disruption [12].

- **Expanding ecosystem:** the presence of platforms, apps, websites, and video channels indicates that learners and schools are already familiar with multiple forms of educational technology [14].

At the same time, e-learning applications in Vietnam face structural limitations that directly motivate Anhoc's design.

- **Access inequality:** home internet access and device ownership vary sharply by income and geography, which means an app cannot assume stable, private, always-on access for every learner [11, 13].
- **Fragmentation of tools and content:** World Bank analysis of Vietnamese e-learning practice notes underdeveloped, unstandardized platforms and inconsistent learning materials, making quality assurance difficult [13].
- **Weak progression support in basic tools:** many digital offerings can deliver lessons or quizzes but do not reliably connect content, assessment history, progression, and administration in one coherent system.
- **Equity risk:** if platform design ignores device constraints, bandwidth limits, or administrative coordination, digital learning can widen rather than reduce educational inequality [11, 13].

These conditions shape the problem statement for Anhoc. A useful platform in Vietnam must not only digitize lessons; it must also provide structured practice, durable assessment records, scalable administration, and a design that remains usable in a context where digital opportunity and digital inequality coexist.

3 Objectives

This thesis has the following objectives:

- document the implemented Anhoc system in a thesis-style format;
- explain the architecture and data model used by the project;
- describe how the platform generates and evaluates mathematical exercises;
- analyze how authentication, authorization, and administration are handled;
- evaluate the strengths, limitations, and next steps of the current implementation.

4 Scope

The scope of this thesis is limited to the code and documentation available in the repository. It focuses on the currently implemented frontend, backend, database schema, and

operational flows. The thesis does not claim to measure learning outcomes through classroom experimentation, because no such experimental dataset is included in the repository. Evaluation is therefore based on software capabilities, system structure, and scenario-level validation.

5 Contributions of the Project

The implemented project contributes several concrete software ideas:

- a bilingual lesson structure for educational content in English and Vietnamese;
- a question-template model that supports variables, derived expressions, constraints, and multiple accepted formulas;
- persisted question snapshots that separate reusable templates from concrete attempt data;
- a progression model combining mastery, XP logs, levels, and achievements;
- a dynamic access-control model that separates actions, resources, and subject restrictions.

6 Thesis Structure

The remainder of this thesis is organized as follows. Chapter 2 presents the requirements and analysis of the project domain. Chapter 3 describes the system design, architecture, security model, and data model. Chapter 4 explains implementation details across the frontend, backend, and database layers. Chapter 5 evaluates the current solution through representative scenarios, strengths, and limitations. Chapter 6 concludes the thesis and outlines future work.

II/ OBJECTIVES

In this chapter, we outline the core functional requirements and design goals of the Anhoc platform. We specify the desired features across user roles and define the expected outcomes of the system implementation.

1 Desired Features

The main goal of the Anhoc project is to implement a structured, adaptive web-based mathematics learning platform with the following services:

- **FEATURE 1: Authentication & Role-Based Access Control**
 - **SUB-FEATURE 1.1: Registration & Login:** Allow new users to create accounts and sign in securely.
 - **SUB-FEATURE 1.2: Dynamic Authorization:** Restrict user actions (create, read, update, delete) based on active roles: Administrator, Supervisor, Teacher, Sub-student, and Free-student.
- **FEATURE 2: Learning & Content Management**
 - **SUB-FEATURE 2.1: Educational Hierarchy:** Organize content by Subjects (e.g., Mathematics), Grades, and Lessons.
 - **SUB-FEATURE 2.2: Bilingual Material:** Support side-by-side study in English and Vietnamese for lessons and practice.
- **FEATURE 3: Dynamic Question Generation & Persisted Snapshots**
 - **SUB-FEATURE 3.1: Question Templates:** Formulate templates with numeric variables, derived calculations, constraints, and multiple accepted answers.
 - **SUB-FEATURE 3.2: Attempt Snapshots:** Persist concrete variables and answers for every attempt to guarantee reproducible grading and review.
- **FEATURE 4: Student Progression & Gamification**
 - **SUB-FEATURE 4.1: Experience Points (XP):** Log XP rewards for practice and test submissions, updating student levels dynamically.
 - **SUB-FEATURE 4.2: Achievements:** Seed and unlock unique badges for completing lessons and scoring well.
- **FEATURE 5: Learn Units & Subscription Management**

- **SUB-FEATURE 5.1: Tenant Isolation:** Group users into Learn Units managed by a Supervisor.
- **SUB-FEATURE 5.2: Subscription Control:** Limit student and teacher slots based on Supervisor plan selections (Free, Pro, Family, Learning Center).
- **FEATURE 6: Interactive Tutoring Chatbot**
 - **SUB-FEATURE 6.1: Step-by-Step AI Guidance:** Embed an AI tutor chatbot capable of parsing mathematics and suggesting solution hints.
 - **SUB-FEATURE 6.2: BYOK Configuration:** Allow users to supply their own API keys to query LLM APIs.

2 Expected Outcome

The implementation aims to construct a production-ready mathematics education portal consisting of:

- A modern Next.js client interface that runs responsively across standard browsers (Chrome, Firefox, Safari) and handles student study, supervisor monitoring, and admin configuration.
- An Express server exposing structured REST endpoints for all data models.
- A PostgreSQL database instance configured with Prisma ORM for type-safe relational schemas.
- A companion Python chatbot microservice implementing conversational math tutoring.

III/ MATERIALS AND METHODS

In this chapter, we present the materials and methods used in the requirements engineering, architectural design, data modeling, and software implementation of the Anhoc platform.

1 Requirements Analysis

1.1 Stakeholders and User Roles

The primary stakeholders in Anhoc are students, administrators, and maintainers of the platform.

- Students consume learning content, complete exercises, take tests, and monitor progress.
- Administrators maintain educational content, templates, users, roles, and reports.
- Maintainers operate the deployment, database, and system configuration.

The codebase also distinguishes between global roles and subject-specific permissions. This is important because a role may have permission to manage some resources without necessarily having access to all subjects.

1.2 Functional Requirements

Table [3.1](#) summarizes the main functional requirements derived from the implementation.

Table 3.1: Core functional requirements

Area	Requirement
Authentication	The system shall allow user registration, login, profile retrieval, password change, and activity retrieval.
Lessons	The system shall support subject and grade organization, lesson listing, lesson detail retrieval, and lesson authoring by authorized users.
Practice	The system shall generate lesson-based practice attempts from reusable templates and store concrete snapshot records for each generated question.
Testing	The system shall create grade-level test attempts, support answer submission, and compute final scores on completion.
Templates	The system shall support creation, update, deletion, preview, reporting, and review of question templates.
Progression	The system shall track study time, mastery status, XP, levels, and achievements.
Administration	The system shall provide user, role, permission, subject-access, and dashboard management endpoints.

1.3 Use Case Model

Figure 3.1 summarizes the main interactions between the implemented role groups and the platform boundary. The repository defines five fixed application roles: **free_student**, **sub_student**, **teacher**, **supervisor**, and **admin**. A guest actor is also useful in the requirements model because registration and login begin before an authenticated role is established.




Figure 3.1: Use case diagram of the Anhoc platform

1.4 Detailed Use Case Diagrams

The following subsections present a detailed use case diagram per functional subsystem, drawn using UML notation. Actors appear outside the system boundary rectangle. Ellipses represent use cases. Solid lines denote associations; dashed lines with open arrowheads denote «include» or «extend» relationships.

Authentication and Account Management




Figure 3.2: Use case diagram — Authentication and Account Management

Learning System




Figure 3.3: Use case diagram — Learning System

Practice and Assessment Engine




Figure 3.4: Use case diagram — Practice and Assessment Engine

AI Math Tutor Chatbot




Figure 3.5: Use case diagram — AI Math Tutor Chatbot

Gamification and Economy




Figure 3.6: Use case diagram — Gamification and Economy

Subscription and Payments




Figure 3.7: Use case diagram — Subscription and Payments

1.5 Use Case Descriptions

Use Case 1: Register or Login

Primary actor	Guest or authenticated platform user
Goal	Create an account or obtain a valid authenticated session.
Preconditions	The user has network access to the frontend and backend services. Login additionally requires an existing account.
Trigger	The user opens the authentication area and submits a registration or login form.

Main flow	(1) The user chooses registration or login. (2) The frontend sends credentials and requested role information to the authentication API. (3) The backend validates the request, resolves the target role, and loads permission data. (4) The backend returns session tokens and profile information. (5) The frontend stores the session and redirects the user to the proper interface.
Postconditions	A new account is created or an authenticated session is established with role-specific permissions.

Use Case 2: Study Lesson

Primary actor	Free student or subscribed student
Goal	Read lesson content for a chosen subject, grade, and lesson.
Preconditions	The student is authenticated and has access to the selected subject.
Trigger	The student opens a lesson from the learning area.
Main flow	(1) The student browses grades and lessons. (2) The frontend requests lesson detail from the API. (3) The backend retrieves lesson metadata and markdown content. (4) The frontend renders the bilingual lesson page. (5) The frontend records study time while the student remains on the page.
Postconditions	The lesson is displayed and study progress can be updated for the student.

Use Case 3: Complete Lesson Practice

Primary actor	Free student or subscribed student
Goal	Practice a lesson through generated questions and receive feedback.
Preconditions	The student is authenticated and the lesson has available question templates.
Trigger	The student starts practice from a lesson page.

Main flow	(1) The frontend requests a practice attempt for the lesson. (2) The backend generates question snapshots from the lesson templates. (3) The student answers questions one by one. (4) The frontend submits responses. (5) The backend evaluates correctness, computes the score, updates mastery and XP, and checks achievements. (6) The result view is returned to the student.
Postconditions	The attempt is stored with snapshot-level grading and the student progression state is updated.

Use Case 4: Take Grade-Level Test

Primary actor	Free student or subscribed student
Goal	Assess broader understanding across the lessons of one grade.
Preconditions	The student is authenticated and the grade has enough templates to assemble a test.
Trigger	The student chooses a grade test from the practice or testing area.
Main flow	(1) The frontend requests a grade-level test attempt. (2) The backend collects eligible templates from lessons in the grade. (3) Question snapshots are generated for the attempt. (4) The student submits answers. (5) The backend finalizes the attempt, calculates the score, and records progression outcomes.
Postconditions	A completed grade-level test attempt is stored and can be reviewed later.

Use Case 5: View Progress and Achievements

Primary actor	Free student or subscribed student
Goal	Review learning progress, XP history, mastery, and earned achievements.
Preconditions	The student is authenticated and has recorded activity in the system.

Trigger	The student opens the dashboard, profile, or achievements page.
Main flow	(1) The frontend requests student statistics and achievement data. (2) The backend loads profile, activity, mastery, and achievement records. (3) The frontend renders summary cards and achievement states.
Postconditions	The student can inspect current progress and earned rewards.

Use Case 6: Manage Content and Question Templates

Primary actor	Teacher, supervisor, or administrator
Goal	Maintain lessons and reusable question templates for the learning platform.
Preconditions	The actor is authenticated and has the required lesson or test management permissions within the permitted subject scope.
Trigger	The actor creates, updates, previews, or deletes a lesson or template in a management interface.
Main flow	(1) The actor edits lesson or template data in the frontend. (2) The frontend sends the request to the backend. (3) The backend validates the JWT, permissions, and subject scope. (4) The backend persists the changes in the database. (5) The frontend refreshes the management view.
Postconditions	The affected content records are updated and become available to the rest of the platform.

Use Case 7: Manage Learning Unit Members and Seats

Primary actor	Supervisor or administrator
Goal	Manage a learning unit by purchasing seats and creating member accounts.
Preconditions	The actor is authenticated, has supervisor-level access, and is associated with a learn unit or is acting as an administrator over one.

Trigger	The actor opens the member management area or seat management flow.
Main flow	(1) The actor requests the current learn unit context. (2) The frontend sends member or seat management operations to the supervisor API. (3) The backend checks whether the user is the relevant supervisor or an administrator. (4) The backend updates seat counts, subscription records, or member accounts. (5) The frontend refreshes the roster and learn unit summary.
Postconditions	Learning unit seats or member accounts are updated consistently with subscription and scope constraints.

Use Case 8: Manage Roles, Access Control, and Dashboard

Primary actor	Administrator
Goal	Control platform-wide users, roles, permissions, subject access, and aggregate reports.
Preconditions	The administrator is authenticated with full management privileges.
Trigger	The administrator opens the admin dashboard or a user and permission management screen.
Main flow	(1) The administrator requests role, user, or dashboard data. (2) The frontend calls the admin API. (3) The backend validates administrative authority. (4) The backend loads or updates users, roles, permissions, and statistical summaries. (5) The frontend renders the resulting dashboard or management state.
Postconditions	Platform-level administration data is viewed or updated with persistent RBAC changes.

1.6 Non-Functional Requirements

The repository implies the following non-functional requirements:

- **Modularity:** frontend, backend, persistence, and progression logic should remain separable.

- **Consistency:** generated questions must remain stable after creation so that grading and reporting are reproducible.
- **Security:** protected endpoints require JWT-based authentication and permission checks.
- **Extensibility:** roles, actions, resources, and subject restrictions should be data-driven rather than hard-coded.
- **Deployability:** the architecture should run in a common cloud arrangement using Vercel, Render, and Neon.

1.7 Learning Content Requirements

Educational content in Anhoc is organized around a hierarchy of subjects, grades, and lessons. Lessons contain bilingual titles and markdown-based content bodies. This leads to several domain requirements:

- one lesson belongs to one grade and one subject;
- lesson content should be readable in more than one language;
- practice content should be attached to lessons through question templates rather than embedded into the lesson body itself;
- progress should be measurable at the lesson level.

1.8 Assessment Requirements

Assessment in the project has two modes:

- lesson-based practice for targeted repetition;
- grade-level tests for broader assessment across available lesson templates.

The assessment subsystem must support question generation, answer submission, correctness evaluation, score calculation, and post-completion progression updates. The implementation also needs to distinguish between theoretical questions and computed math questions because the grading strategy differs between those two cases.

1.9 Administrative Requirements

Administrative functionality in Anhoc goes beyond simple content editing. The platform must manage:

- users and their assigned roles;

- permissions by action and resource;
- subject-level restrictions per role;
- dashboard statistics and user insights;
- reports submitted for problematic generated questions.

This requirement set justifies the presence of a dynamic RBAC model instead of a fixed boolean admin flag.

1.10 Operational Assumptions

The codebase assumes a web-first deployment model:

- the frontend is delivered as a Next.js application;
- the backend exposes REST endpoints under `/api/v1`;
- PostgreSQL is the source of truth for transactional and analytical learning data;
- environment variables define API URLs, CORS origins, JWT secret material, and achievement seeding behavior.

These assumptions shape the rest of the architecture described in the next chapter.

2 System Design

2.1 Architectural Overview

Anhoc uses a three-tier web architecture. A Next.js frontend built on React provides the user interface. An Express backend exposes REST endpoints and business logic. PostgreSQL stores content, user identities, generated snapshots, and progression data. Prisma acts as the ORM boundary between the backend and the database [1, 2, 3, 4, 5].




Figure 3.8: Architecture design of the Anhoc platform

At a high level, the request flow is:

1. the browser interacts with the Next.js frontend;
2. the frontend calls backend routes under **/api/v1**;
3. the backend validates identity and permissions;
4. business services coordinate data access and progression updates;
5. PostgreSQL persists authoritative state.

2.2 Deployment Design

The repository documents a practical deployment split:

- Vercel hosts the frontend runtime;
- Render hosts the Node.js API;
- Neon provides PostgreSQL infrastructure.

This split is appropriate for a full-stack TypeScript project because it separates the static or server-rendered frontend from the API runtime and keeps the relational database independent of either application tier.

2.3 Backend Design

The backend entry point is `backend/server.ts`. It initializes Express, configures CORS, enables JSON parsing, mounts route modules, exposes a health endpoint, and conditionally seeds achievements at startup. The route structure is organized by domain:

- `auth.ts` for identity, login, and profile operations;
- `lessons.ts` for educational content and practice initiation;
- `tests.ts` for attempts, answer submission, template management, and reports;
- `admin.ts` for users, roles, access control, and statistics;
- `achievements.ts` for achievement retrieval and re-evaluation.

This separation keeps HTTP concerns close to route boundaries while pushing calculation-heavy logic into services.

2.4 Frontend Design

The frontend is built with Next.js using an app-router structure under `frontend/src/app`. It includes separate areas for authentication, student flows, and administration. The frontend is supported by Axios-based service modules, Redux state for authentication and permissions, and feature-specific components such as lesson cards, question players, result modals, and achievement galleries. The repository also uses `next-intl` for localized dictionaries and request-scoped internationalization in the App Router [9].

From a design perspective, the frontend acts as an orchestration layer:

- it collects user input and triggers backend operations;
- it renders lesson, practice, test, and admin views;
- it persists authentication tokens in browser storage and cookies;
- it enforces route-level and component-level access decisions using the permission payload returned by the backend.

2.5 Security Model

Authentication is handled with JWT bearer tokens. The middleware in `backend/middleware/auth.ts` extracts the token from the `Authorization` header, verifies it using the configured secret, and attaches the decoded payload to the request object [6, 7].

Authorization is layered on top of authentication. The project uses:

- role identity for coarse-grained decisions;
- grouped resource permissions for action checks;
- subject-level role restrictions for scoped access;
- helper middleware such as `isAdmin` and `selfOrAdmin`.

The design choice to keep permissions data-driven is significant. It allows new resources and actions to be introduced through database records and route logic instead of forcing the codebase into a rigid two-role system.

2.6 Data Model

The database schema in `backend/prisma/schema.prisma` defines the educational, assessment, and progression domain. Table 3.10 summarizes the most important entities.

Table 3.10: Main persistent entities

Entity	Purpose
<code>users</code>	Stores account identity and links to roles, attempts, mastery, achievements, and statistics.
<code>roles</code> , <code>actions</code> , <code>resources</code> , <code>role_permissions</code>	Implements dynamic role-based access control.
<code>subjects</code> , <code>grades</code> , <code>lessons</code>	Represents the educational content hierarchy.
<code>question_templates</code>	Stores reusable question definitions and logic.
<code>test_attempts</code>	Stores headers for practice and test sessions.
<code>question_snapshots</code>	Stores concrete generated questions bound to attempts.
<code>user_lesson_mastery</code>	Stores lesson-level progress and completion status.
<code>student_stats</code> , <code>xp_logs</code>	Stores aggregate progression state and XP history.
<code>achievements</code> , <code>user_achievements</code>	Stores gamification rewards and earned records.
<code>question_template_reports</code>	Stores reports for problematic questions.
<code>audit_logs</code> , <code>activity_evidence</code>	Supports traceability and evidence collection.

2.7 Entity-Relationship Diagrams

The following diagrams use a simplified ERD notation. Each entity is a rectangle with its name in the header row and its key attributes listed below. Lines between entities carry crow's foot cardinality markers: $||$ for exactly one, $o|$ for zero-or-one, $o\{$ for zero-or-many, $|\{$ for one-or-many. Primary keys are marked **PK**; foreign keys are marked **FK**.

ERD Group 1 — User and Authentication




Figure 3.9: ERD Group 1 — User & Authentication tables

This group contains the core identity and access control entities. Every other table in the schema references **users** as its ultimate root.

Table: users The central entity of the entire platform. Each row represents one registered account.

Column	Type	Description
id	UUID (PK)	Auto-generated primary key using <code>gen_random_uuid()</code> .
username	VARCHAR(50)	Unique display name chosen at registration.
email	VARCHAR(255)	Unique email address; used for login and notifications.
password_hash	TEXT	bcrypt hash of the user's password. Never stored in plaintext.
account_status	VARCHAR(20)	Lifecycle state: active , pending , suspended , or deleted .
role_id	INT (FK)	References roles.id ; determines permissions via RBAC.
learn_unit_id	UUID (FK, null)	References learn_units.id ; set for group-plan members.
stripe_customer_id	TEXT (unique, null)	Stripe customer identifier; set on first checkout.
active_session_id	TEXT (null)	UUID of the current valid session; rotated on every login for single-session enforcement.
avatar_url	VARCHAR(255)	URL of the active equipped avatar image.
first_name, last_name	VARCHAR(100)	Optional display name components.
created_at	TIMESTAMPTZ	Account creation timestamp.

Table: user_securities One-to-one extension of **users** holding all sensitive security metadata. Separated from the main users table to avoid loading security data on every profile read.

Column	Type	Description
user_id	UUID (PK/FK)	References users.id with CASCADE delete.
email_verification_token	VARCHAR(255)	Time-limited signed token sent in the activation email.

email_verified_at	TIMESTAMPZ	Timestamp of successful email verification; null if unverified.
failed_login_attempts	INT	Counter reset to zero on successful login.
lockout_until	TIMESTAMPZ	If set, login is blocked until this timestamp (account lockout).
password_reset_token	VARCHAR(255)	Signed reset token; cleared after use or expiry.
password_reset_sent_at	TIMESTAMPZ	Throttle timestamp to prevent reset email flooding.

Table: refresh_tokens Stores active refresh tokens issued during login. A user may have multiple rows (e.g. across devices), but single-session enforcement via `active_session_id` invalidates all but the latest.

Column	Type	Description
id	UUID (PK)	Row identifier.
token	TEXT (UNIQUE)	The opaque refresh token string stored in the <code>HttpOnly</code> cookie.
user_id	UUID (FK)	Owner of this token. CASCADE deleted with the user.
expires_at	TIMESTAMPZ	Hard expiry; token is rejected after this even if not revoked.
revoked_at	TIMESTAMPZ (null)	Set when the token is explicitly invalidated (logout or rotation).

Table: oauth_accounts Stores OAuth2 identity links for social login providers (Google, etc.). One user may link multiple providers.

Column	Type	Description
id	UUID (PK)	Row identifier.
user_id	UUID (FK)	References <code>users.id</code> .
provider	VARCHAR(50)	Provider name, e.g. <code>google</code> , <code>facebook</code> .
provider_user_id	VARCHAR(255)	The unique user ID within the external provider's system.

Table: roles and role_permissions Implement the Role-Based Access Control (RBAC) model. Each role has a set of (`action`, `resource`) permission pairs. Mid-

dleware checks these at runtime for every protected route.

ERD Group 2 — Curriculum and Content




Figure 3.10: ERD Group 2 — Curriculum & Content tables

This group models the three-level curriculum hierarchy: Subject → Grade → Lesson → Question Template.

Table: subjects Top-level subject classification (e.g. Mathematics, Physics).

Column	Type	Description
id	INT (PK)	Auto-increment integer key.
slug	VARCHAR(50)	Unique display name chosen at registration.
title_en / title_vi	VARCHAR(100)	Bilingual display names.
is_classified	BOOLEAN	If true, access requires an approved <code>user_subject_access_request</code> .

color	VARCHAR(20)	Hex colour for UI theming of this subject.
created_by	UUID (FK, null)	Administrator who created the subject.

Table: grades Represents a school year/grade within a subject (e.g. “Grade 6 Mathematics”).

Column	Type	Description
id	INT (PK)	Auto-increment key.
slug	VARCHAR(50)	URL-safe grade identifier, e.g. grade-6 .
subject_id	INT (FK, null)	Parent subject.
title_en / title_vi	VARCHAR(100)	Bilingual grade titles.

Table: lessons The core content entity. Each lesson belongs to one grade and one subject, and contains full bilingual theory text in Markdown format.

Column	Type	Description
id	UUID (PK)	Row identifier.
grade_id	INT (FK)	Parent grade.
subject_id	INT (FK)	Parent subject (denormalised for fast filtering).
title_en / title_vi	VARCHAR(255)	Bilingual lesson titles, displayed on the roadmap.
content_markdown_en / vi	TEXT	Full theory content in Markdown with embedded KaTeX.
order_index	INT (null)	Display ordering within the grade roadmap.
is_premium	BOOLEAN	If true, only subscribers can view the full content.
updated_at	TIMESTAMPTZ	Last content update; used to invalidate Redis caches.

Table: question_templates Parameterised question blueprints. At practice time, numeric parameters are instantiated from the `logic_config` ranges and the ground-truth answer is computed by SymPy from `accepted_formulas`.

Column	Type	Description
--------	------	-------------

id	UUID (PK)	Row identifier.
lesson_id	UUID (FK, null)	Associated lesson; null for topic-level free questions.
template_type	VARCHAR(50)	Category, e.g. <code>multiple_choice</code> , <code>open_input</code> .
difficulty	VARCHAR(20)	One of: <code>easy</code> , <code>medium</code> , <code>hard</code> .
body_template_en / vi	TEXT	Question text with <code>{{a}}</code> , <code>{{b}}</code> parameter placeholders.
logic_config	JSON	Parameter ranges and types, e.g. <code>{"a": [2,20], "b": [1,9]}</code> .
accepted_formulas	TEXT[]	SymPy expression strings to compute the correct answer.
tags	TEXT[]	Topic tags for filtering and adaptive selection.

Table: user_subject_access_requests Students request access to classified subjects. Administrators approve or reject via the admin panel.

ERD Group 3 — Practice, Tests and Progress




Figure 3.11: ERD Group 3 — Practice, Tests & Progress tables

This group maintains records of active studies, completed lessons, test score achievements, and audit logs.

Table: test_attempts Records every practice session or grade test attempt. The `is_practice` flag distinguishes practice runs from formal tests.

Column	Type	Description
--------	------	-------------

id	UUID (PK)	Attempt identifier.
user_id	UUID (FK, null)	Null for anonymous guest game attempts.
lesson_id	UUID (FK, null)	Lesson context; null for grade-level tests.
total_score	DECIMAL(5,2)	Percentage score, computed on completion.
is_practice	BOOLEAN	True for lesson practice; false for grade tests.
is_completed	BOOLEAN	Set to true when the student submits or the timer expires.
started_at	TIMESTAMPTZ	Session start time (used for speed calculation).
completed_at	TIMESTAMPTZ	Session end time. Duration = <code>completed_at - started_at</code> .

Table: question_snapshots One row per question per attempt. Captures the instantiated parameter values, the student’s answer, and the ground-truth answers at the time of the attempt (immutable record for audit purposes).

Column	Type	Description
id	UUID (PK)	Snapshot identifier.
attempt_id	UUID (FK, null)	Parent attempt. CASCADE deleted.
template_id	UUID (FK, null)	Source template (SET NULL on template deletion for history).
generated_variables	JSON	E.g. <code>{"a": 7, "b": 3}</code> — the random values used for this instance.
student_answer	TEXT (null)	The answer submitted by the student.
right_answers	TEXT[]	Array of accepted correct answer strings at generation time.
is_correct	BOOLEAN	Whether <code>student_answer</code> matches any element of <code>right_answers</code> .
points_earned	INT	Points scored (0 or configured per-template value).

Table: user_lesson_mastery Composite-primary-key junction tracking each student’s cumulative performance per lesson. Updated atomically on every attempt completion.

Column	Type	Description
user_id	UUID (PK/FK)	Composite PK component. References <code>users.id</code> .
lesson_id	UUID (PK/FK)	Composite PK component. References <code>lessons.id</code> .
mastery_score	DECIMAL(5,2)	Rolling weighted-average performance score (0–100).
total_study_time	INT	Cumulative seconds spent on the lesson's reading sessions.
total_test_time	INT	Cumulative seconds spent on this lesson's practice attempts.
completion_status	VARCHAR	One of: <code>not_started</code> , <code>in_progress</code> , <code>completed</code> . Drives the roadmap badge.

Table: `learning_sessions` and `question_template_reports` `learning_sessions` stores each discrete reading duration to measure study metrics. `question_template_reports` lets students flag issues with question accuracy, linking templates, snapshots, and logs for review.

ERD Group 4 — Gamification and Economy




Figure 3.12: ERD Group 4 — Gamification & Economy tables

This group maintains points, levels, items, streak configurations, daily quest targets, and earned achievements.

Table: `student_stats` One-to-one extension of `users` holding the student’s entire gamification state. Kept separate from `users` to avoid locking the user row on every XP award transaction.

Column	Type	Description
<code>user_id</code>	UUID (PK/FK)	One-to-one with <code>users</code> .
<code>total_xp</code>	INT	Cumulative experience points earned.
<code>level</code>	INT	Current level, derived from <code>total_xp</code> against the levelling formula.
<code>average_score</code>	DECIMAL(5,2)	Rolling average across all practice attempts.
<code>lives</code>	INT	Current number of lives (1 consumed per practice attempt).
<code>last_life_restored_at</code>	TIMESTAMPTZ	Timestamp of last hourly life restore; drives the restore scheduler.
<code>coins</code>	INT	Current Ancoin balance.
<code>level_points</code>	INT	Unspent attribute points earned on level-up.
<code>inventory</code>	JSON	Map of owned shop item holdings.
<code>equipped_items</code>	JSON	Map of currently active cosmetic items.
<code>upgrades</code>	JSON	Map of purchased permanent upgrades.

Table: `xp_logs` Immutable append-only log of every XP award event. Used for the XP history view and for audit purposes.

Column	Type	Description
<code>user_id</code>	UUID (FK)	The student who received the XP.
<code>amount</code>	INT	XP delta awarded.
<code>reason</code>	VARCHAR(100)	Human-readable label (e.g. <code>practice_complete</code>).
<code>occurred_at</code>	TIMESTAMPTZ	Timestamp of award event.

Table: `user_streaks`, `achievements`, `themes`, `user_daily_quests` Tracks continuous engagement records, custom variable sets for light/dark CSS styles, badges

catalogued by milestones, and recurring daily quest logs designed to reward incremental student learning.

ERD Group 5 — Subscriptions, Gaming and Notifications




Figure 3.13: ERD Group 5 — Subscriptions, Gaming and Notifications

This group governs subscription limits, Stripe transactions, challenges/attempts in multiplayer gaming, and system-wide notifications.

Table: plans Master catalogue of subscription plans (e.g. Free, Pro, Family). Stores base pricing in both monthly and annual cycles, plus per-unit pricing multipliers for flexible enterprise pricing.

Column	Type	Description
id	INT (PK)	Auto-increment plan identifier.
name	VARCHAR(50)	Internal slug: <code>free</code> , <code>pro</code> , <code>family</code> , <code>learning_center</code> .
vi_name / en_name	VARCHAR(100)	Bilingual marketing display names.
price_monthly / annually	DECIMAL(10,2)	Base price per billing cycle.
max_students / teachers / lessons	INT (null)	Feature usage caps; null means unlimited.
effect_from / effect_to	TIMESTAMPTZ	Validity window enabling time-based pricing changes without deleting old plans.

Table: subscription_plans Each row is one active or historical subscription instance for a user. The `stripe_subscription_id` links to the Stripe subscription object for synchronisation via webhooks.

Column	Type	Description
id	UUID (PK)	Subscription instance identifier.
user_id	UUID (FK)	Subscriber.
plan_id	INT (FK)	References <code>plans.id</code> .
status	VARCHAR(20)	One of: <code>active</code> , <code>cancelled</code> , <code>expired</code> , <code>superseded</code> .
billing_cycle	VARCHAR(20)	<code>monthly</code> or <code>annually</code> .
auto_renew	BOOLEAN	Whether Stripe should attempt renewal at period end.
stripe_subscription_id	TEXT (UNIQUE)	Stripe object ID, used for webhook upsert matching.
max_students / lessons / grades	INT (null)	Custom limits captured at check-out time for per-seat enterprise plans.

Table: game_challenges A competitive battle session created by a student. The `code` field is a short shareable room code. The `config` JSON stores the pre-generated question set for the challenge to ensure all participants answer the same questions.

Column	Type	Description
id	UUID (PK)	Challenge identifier.
code	VARCHAR(10) (UNIQUE)	Short room code shared with opponents.
game_type	VARCHAR(50)	Battle mode variant: speed , climb , match .
grade_id / lesson_id	FK (null)	Scope of the challenge questions.
created_by	UUID (FK)	Creator; also the host.
is_active	BOOLEAN	Set to false when the challenge ends.
config	JSON	Frozen question set for this challenge instance.

Table: invoices, game_attempts, notifications, audit_logs Stores Stripe transactions, battle participation logs (including guests), polymorphic system-wide alerts, and immutable logs detailing sensitive admin actions or critical events.

Schema Summary

Table 3.27: Complete Database Schema — Table Summary (38 tables)

Table	Group	Purpose
users	Auth	Core user account entity
user_securities	Auth	Security tokens, logout state
refresh_tokens	Auth	JWT refresh token store
oauth_accounts	Auth	Social login provider links
roles	Auth	RBAC role definitions
actions / resources	Auth	RBAC action and resource catalogue
role_permissions	Auth	Role $\times$ Action $\times$ Resource permission matrix
learn_units	Auth	Group subscription unit (supervisor + members)
subjects	Curriculum	Subject top-level catalogue
grades	Curriculum	Grade/year level within a subject
lessons	Curriculum	Theory content pages with bilingual Markdown
question_templates	Curriculum	Parameterised question blueprints

user_subject_access_requests	Curriculum	Access approval workflow for classified subjects
role_subject_permissions	Curriculum	Subject-level permission overrides per role
user_learning_profiles	Curriculum	Preferred subject stored per student
test_attempts	Practice	Practice or test session record
question_snapshots	Practice	Per-question instantiated attempt record
user_lesson_mastery	Practice	Cumulative per-lesson performance state
learning_sessions	Practice	Reading session time log
activity_evidence	Practice	Uploaded evidence images per attempt
question_template_reports	Practice	Student-flagged question error reports
student_stats	Gamification	XP, level, lives, coins, inventory, upgrades
xp_logs	Gamification	Immutable XP event ledger
achievements	Gamification	Achievement definition catalogue
user_achievements	Gamification	Student $\times$ Achievement unlock junction
themes	Gamification	CSS variable theme definitions
user_streaks	Gamification	Per-student streak state
streak_rewards	Gamification	Admin-configured milestone reward definitions
user_streak_reward_claims	Gamification	Anti-double-claim claim ledger
user_daily_quests	Gamification	Daily generated quests per student
user_follows	Social	Student follow relationships
notifications	System	Polymorphic notification store
audit_logs	System	Append-only system event log
plans	Subscription	Subscription plan catalogue
subscription_plans	Subscription	Per-user subscription instances
invoices	Subscription	Billing event records
game_challenges	Gaming	Battle session definitions
game_attempts	Gaming	Per-player battle participation records

2.8 Question Snapshot Strategy

One of the strongest design decisions in Anhoc is the separation between `question_templates` and `question_snapshots`. Templates describe how a question can be gen-

erated. Snapshots record the exact generated variables, accepted answers, student response, correctness result, and awarded points for one attempt.

This decision solves several engineering problems:

- a generated question does not change after the student starts an attempt;
- grading can be reproduced later without regenerating random variables;
- question reporting can refer to a concrete faulty instance rather than an abstract template;
- analytics can distinguish between template quality and attempt outcomes.

2.9 Progression Design

Progression in Anhoc is not limited to a final score. The backend updates several layers of state:

- lesson mastery based on lesson-linked practice results;
- XP logs for additive progression history;
- aggregate student statistics for dashboards and profile displays;
- achievements that can also grant extra XP.

This design supports both immediate feedback and longitudinal tracking, which is more appropriate for learning software than a simple pass or fail test history.

2.10 Sequence Diagrams for Main Use Cases

The following sequence diagrams refine the use cases introduced in the requirements chapter and map them to the implemented frontend, backend, service, and database layers.

Use case 1: Register or Login




Figure 3.14: Sequence diagram for registration or login

Use case 2: Study Lesson




Figure 3.15: Sequence diagram for studying a lesson

Use case 3: Complete Lesson Practice




Figure 3.16: Sequence diagram for completing lesson practice

Use case 4: Take Grade-Level Test




Figure 3.17: Sequence diagram for taking a grade-level test

Use case 5: View Progress and Achievements




Figure 3.18: Sequence diagram for viewing progress and achievements

Use case 6: Manage Content and Question Templates




Figure 3.19: Sequence diagram for managing content and question templates

Use case 7: Manage Learning Unit Members and Seats




Figure 3.20: Sequence diagram for managing learning unit members and seats

Use case 8: Manage Roles, Access Control, and Dashboard




Figure 3.21: Sequence diagram for managing roles, access control, and dashboard data

3 System Implementation

3.1 Frontend Implementation

The frontend is implemented with Next.js and TypeScript. The app-router structure separates student and administrative workflows into route segments, while service modules in `frontend/src/services` centralize communication with the backend.

Important frontend concerns include:

- authentication bootstrap and persistence;
- protected routes and permission checks;
- lesson and practice page composition;
- test flows with answer submission and result presentation;
- admin screens for users, roles, reports, and dashboards.

The component structure indicates a feature-oriented approach rather than a single monolithic page design. Examples include `QuestionPlayer.tsx`, `PracticeResultModal.tsx`, `AchievementGallery.tsx`, and `CreateTemplateModal.tsx`.

3.2 Backend API Implementation

The backend exposes versioned REST endpoints. Table 3.28 groups the major API areas.

Table 3.28: Main API groups

API Group	Purpose
<code>/api/v1/auth</code>	Registration, login, profile retrieval, password change, and activity endpoints.
<code>/api/v1/lessons</code>	Lesson, grade, subject, mastery, practice initiation, and study-time endpoints.
<code>/api/v1/tests</code>	Grade tests, attempt detail, answer submission, finish flow, template management, and report workflow.
<code>/api/v1/admin</code>	Users, roles, access control, insights, and dashboard statistics.
<code>/api/v1/achievements</code>	Achievement seeding, listing, earned-state retrieval, and re-checking.

The API design is pragmatic. It is expressive enough for the frontend and explicit enough for debugging, while remaining simple to secure and deploy.

3.3 Question Generation and Grading

The file `backend/services/mathService.ts` contains a central part of the project: logic-driven variable generation and answer evaluation. The implementation uses the `mathjs` expression engine for parsing and evaluating formulas, derived expressions, and constraints [8]. The service supports several rule types:

- fixed numeric values;
- arrays of possible values;
- numeric ranges with optional step and precision;
- random choices from configured lists;
- derived expressions evaluated from previously generated variables;
- constraints that must hold before a question is accepted.

The generation flow can be summarized as:

1. parse the stored `logic_config`;
2. generate variables from configured rules;
3. apply derived expressions;
4. verify constraints;
5. repeat until a valid assignment is found or a maximum number of attempts is reached.

Answer checking supports multiple cases:

- numeric answers, compared with a small tolerance;
- boolean answers, normalized from common textual forms;
- text or structured answers, normalized for comparison;
- multiple accepted formulas for the same question.

This is a stronger implementation than a naive string-equality checker because it tolerates common input variance while still grounding correctness in mathematically evaluated formulas.

3.4 Attempt Lifecycle

The implemented attempt lifecycle is one of the most important interactions in the system.

Practice Start

When a student starts practice for a lesson:

1. the backend validates lesson existence and subject access;
2. it loads lesson templates;
3. it generates variables and right answers for selected templates;
4. it creates a `test_attempt` record;
5. it bulk-inserts `question_snapshots`;
6. it returns a stable attempt payload to the frontend.

Answer Submission

When a student submits an answer:

1. the backend loads the snapshot and template;
2. it evaluates the answer using either text comparison or formula-based checking;
3. it updates the stored snapshot with answer and correctness state;
4. it returns correctness information and explanation content.

Attempt Completion

When an attempt is finished:

1. unanswered snapshots are marked incorrect;
2. correct answers are counted and the final score is computed;
3. lesson mastery is updated when the attempt is linked to a lesson;
4. XP is awarded and logged;
5. aggregate student statistics are recalculated;
6. achievements are re-evaluated and may award additional XP.

This completion pipeline shows that the project treats scoring as part of a broader progression workflow rather than as an isolated calculation.

3.5 Administrative Implementation

The administrative subsystem includes user and role management, subject permission management, dashboard statistics, and question report review. The route design

suggests that administrators can change both coarse and fine-grained access behavior without redeploying the application.

This is reinforced by the underlying data model:

- `roles` store named roles;
- `actions` and `resources` define the authorization vocabulary;
- `role_permissions` map actions to resources per role;
- `role_subject_permissions` restrict which subjects a role can manage.

This structure scales better than a hard-coded permission matrix embedded directly in frontend components.

3.6 Configuration and Deployment

The repository defines environment variables for both frontend and backend execution. Examples include:

- API base URLs and rewrite targets for the frontend;
- `DATABASE_URL`, `JWT_SECRET`, and `CORS_ORIGINS` for the backend;
- `AUTO_SEED_ACHIEVEMENTS` to control whether achievement definitions are seeded automatically.

The backend startup process explicitly checks the achievement seeding flag before inserting seed data. This is a sound operational choice because it prevents accidental reseeding in every environment.

To support standard containerised hosting on services like Render, the backend includes a custom `Dockerfile` built on the lightweight `node:22-alpine` base image. The configuration installs runtime dependencies such as `openssl` (required for the Prisma Client runtime) and `curl`. A container health check is declared natively using a Node.js fetch script targeting the backend health endpoint (`/api/v1/health`) on the configured port.

A key operational challenge resolved in the deployment configuration involves build-time database constraints. During the Docker compilation phase on Render, the `DATABASE_URL` environment variable is not injected (as it is scoped strictly to the application's runtime phase). The compiler initially failed when executing `npxprisma generate` because the local ORM configuration (`prisma.config.ts`) tried to resolve this variable dynamically using Prisma's `env()` configuration helper, which enforces lookup checks during build time. This constraint was resolved by refactoring the configuration file to reference the standard Node.js `process.env.DATABASE_URL` lookup directly. This

adjustment allows the client generation script to execute successfully during image compilation, decoupling the database connectivity verification from the build phase.

3.7 User Interface Presentation

The frontend interface features a responsive layout designed for both desktop side navigation and mobile layout constraints. For small screen widths, a floating navigation bar (`BottomNav.tsx`) provides touch-friendly route switching between student dashboard, roadmap, practice sessions, chatbot tutor, and settings. To enhance visual separation and ensure the bottom navigation stands out over the dynamic scrolling content, the element implements absolute layout positioning with a top border class (`border-t-2`) and dynamic slide-in entry animation.

IV/ RESULTS AND DISCUSSION

1 Method of Evaluation

This thesis evaluates Anhoc through implementation analysis and scenario-based validation rather than controlled educational experiments. The repository does not include classroom outcome data, benchmark reports, or a formal automated test suite that would support statistical claims. It would be incorrect to invent such evidence. Instead, the evaluation focuses on whether the implemented design addresses the stated software goals.

2 Representative User Scenarios

Scenario 1: Student Login and Study Access

The system supports a full login flow in which credentials are submitted from the frontend, verified by the backend, and returned as a JWT with grouped permissions. This meets the requirement for authenticated access to protected student and admin workflows.

Scenario 2: Lesson-Based Practice

The platform can start practice sessions from lesson templates, generate concrete question instances, and return a stable attempt payload. Because snapshots are persisted, a student can answer the exact generated questions rather than a regenerated variant. This directly satisfies the consistency requirement discussed earlier.

Scenario 3: Assessment and Progression

The finish-attempt flow does more than assign a score. It updates lesson mastery, adds XP, recalculates student statistics, and checks achievements. This makes the system better aligned with long-term learning engagement than a basic quiz engine.

Scenario 4: Administrative Control

The platform exposes dedicated endpoints for roles, users, access control, reports, and dashboard statistics. That breadth indicates the system was designed as a manageable

platform rather than a single-user prototype.

3 Observed Strengths

Several strengths are visible in the current implementation:

- **Clear layering.** Frontend, backend, services, and database concerns are separated cleanly.
- **Snapshot persistence.** The snapshot model is the strongest architectural decision in the project and improves grading stability, traceability, and debugging.
- **Flexible authorization.** Dynamic RBAC plus subject-level permissions provides useful operational flexibility.
- **Educational progression.** Mastery, XP, levels, and achievements create a richer student model than raw scores alone.
- **Template expressiveness.** Variable ranges, derived expressions, constraints, and accepted formulas make the question system meaningfully reusable.

4 Current Limitations

The project also has clear limitations that should be acknowledged directly.

- **No formal experimental validation.** The repository does not provide evidence of student learning impact.
- **Limited compile-time guarantees across dynamic JSON logic.** Template logic is flexible, but JSON-configured expressions can still fail at runtime.
- **Potential duplication of schema ownership.** The repository contains Prisma schema material in both frontend and backend directories, which creates a risk of divergence if not managed carefully.
- **Operational maturity is still emerging.** There is little evidence in the repository of broad automated testing, observability, or release-hardening practices.
- **Reporting and analytics can be expanded.** The current design supports insights and statistics, but a deeper learning analytics layer is not yet visible.

5 Risks and Tradeoffs

The design choices in Anhoc involve reasonable tradeoffs:

- Persisting snapshots increases storage cost, but the gain in auditability and reproducibility is worth it.
- Dynamic permissions improve flexibility, but they also require careful synchronization between backend enforcement and frontend affordances.
- Formula-driven generation increases content reusability, but malformed expressions can be harder to validate than fixed question banks.

6 Future Improvements

Several next steps would materially improve the project:

- add automated tests for route behavior, template validation, and progression logic;
- build a stronger authoring validation layer for `logic_config` and accepted formulas;
- add richer analytics for question difficulty, error patterns, and lesson effectiveness;
- improve observability through structured logs, metrics, and failure dashboards;
- define a single canonical Prisma schema source and remove duplication risk;
- conduct user studies or classroom pilots to validate educational effectiveness.

7 Discussion

Taken as a software engineering project, Anhoc is already beyond a simple tutorial application. It contains meaningful domain modeling, a non-trivial assessment engine, a flexible authorization model, and a progression subsystem tied to learning behavior. Its strongest academic value lies in how it combines educational workflows with practical backend architecture decisions. The platform would benefit from deeper validation and more operational hardening, but the core design is coherent and defensible.

V/ CONCLUSION & PERSPECTIVE

This thesis documented the design and implementation of Anhoc, a web-based mathematics learning platform that combines bilingual lesson content, generated practice exercises, grade-level assessments, progression tracking, and administrative control. The implemented system uses a modern web stack with a Next.js frontend, an Express backend, Prisma as the persistence layer, and PostgreSQL as the database.

The project demonstrates several sound software engineering decisions. Most importantly, it separates reusable question templates from persisted question snapshots, enabling reproducible grading and better issue reporting. It also uses a dynamic access-control model and a modular progression pipeline for mastery, XP, and achievements. These choices make the system more extensible than a minimal educational CRUD application.

At the same time, the current implementation remains an evolving platform rather than a finished production learning system. It still needs broader automated testing, stronger validation of author-authored template logic, clearer schema ownership, and empirical studies of educational effectiveness. Even with these limitations, the project is a solid foundation for future academic work and practical product development in mathematics learning software.

APPENDICES

A.1 Build and Run Notes

The repository documents separate frontend and backend setup processes. A representative local workflow is shown below.

```
cd backend
npm install
npm run prisma:migrate
npm run dev
```

```
cd frontend
npm install
npm run dev
```

A.2 Selected Environment Variables

Table A.1: Representative environment variables

Variable	Purpose
NEXT_PUBLIC_API_URL	Browser-visible API base URL for the frontend.
INTERNAL_API_URL	Server-side fetch base used by the frontend runtime.
BACKEND_URL	Optional rewrite target for same-origin frontend calls.
SERVER_PORT	Backend listening port.
DATABASE_URL	PostgreSQL connection string for Prisma.
JWT_SECRET	Secret used to verify and sign JWT tokens.
CORS_ORIGINS	Allowed origins for cross-origin backend requests.
AUTO_SEED_ACHIEVEMENTS	Controls automatic achievement seeding on startup.

A.3 Selected REST Endpoints

Table A.2: Example endpoints from the implemented API

Method	Path	Purpose
POST	/api/v1/auth/login	Authenticate a user and return token plus permission payload.
POST	/api/v1/lessons/:id/practice	Start a lesson-based practice attempt.
POST	/api/v1/tests/submit-answer	Submit an answer for a generated question snapshot.
POST	/api/v1/tests/attempts/:id/finish	Complete an attempt and trigger progression updates.
GET	/api/v1/admin/access-control	Retrieve data needed for role and permission administration.

A.4 Application User Interface Screenshots

The following pages illustrate the implemented user interfaces of the Anhoc platform, captured during live local execution of both student and administrative portals.




Figure A.1: Student dashboard displaying experience points, active level, and unlocked achievements.




Figure A.2: Grade-level learning pathway showing available lessons and bilingual subject content.




Figure A.3: Interactive practice interface featuring a dynamically generated mathematics question and live evaluation feedback.




Figure A.4: Achievements gallery displaying seeded badges and their associated experience point rewards.




Figure A.5: Tutoring chatbot widget answering student queries with step-by-step mathematical reasoning.



Figure A.6: Administrative overview panel showing system stats, active users, and lesson metrics.




Figure A.7: Admin user management dashboard showing all registered users and their system roles.




Figure A.8: Access control settings panel for configuring user seats and editing role-based permission profiles.

BIBLIOGRAPHY

- [1] Vercel. *Next.js Documentation: App Router*. Available at: <https://nextjs.org/docs/app>. Accessed June 15, 2026.
- [2] Meta Platforms, Inc. and contributors. *React Documentation*. Available at: <https://react.dev/learn>. Accessed June 15, 2026.
- [3] OpenJS Foundation and Express contributors. *Express.js 5.x API Reference*. Available at: <https://expressjs.com/en/api/>. Accessed June 15, 2026.
- [4] Prisma. *Prisma ORM Documentation*. Available at: <https://www.prisma.io/docs/orm>. Accessed June 15, 2026.
- [5] PostgreSQL Global Development Group. *PostgreSQL 18 Documentation*. Available at: <https://www.postgresql.org/docs/current/index.html>. Accessed June 15, 2026.
- [6] M. Jones, J. Bradley, and N. Sakimura. *RFC 7519: JSON Web Token (JWT)*. RFC Editor, May 2015. Available at: <https://www.rfc-editor.org/info/rfc7519/>.
- [7] N. Sheffer, D. Hardt, and M. Jones. *RFC 8725: JSON Web Token Best Current Practices*. RFC Editor, February 2020. Available at: <https://www.rfc-editor.org/info/rfc8725/>.
- [8] *mathjs Documentation*. Available at: <https://mathjs.org/docs/>. Accessed June 15, 2026.
- [9] Jan Amann and contributors. *next-intl Documentation: Getting Started*. Available at: <https://next-intl.dev/docs/getting-started>. Accessed June 15, 2026.
- [10] Kepios, DataReportal, and We Are Social. *Digital 2024: Vietnam*. Available at: <https://datareportal.com/reports/digital-2024-vietnam>. Accessed June 22, 2026.
- [11] UNESCO Global Education Monitoring Report. *UNESCO and SEAMEO Report: Technology Revolutionizing Education in Viet Nam and the Region, but Not a One-Size-Fits-All Solution*. Available at: <https://www.unesco.org/gem-report/en/articles/unesco-and-seameo-report-technology-revolutionizing-education-viet-nam-and-r>. Accessed June 22, 2026.

- [12] World Bank. *Viet Nam: Hoc Tren Truyen Hinh (Distance Learning through TV Broadcasting)*. Available at: <https://documents1.worldbank.org/curated/en/848771599124103455/pdf/Vietnam-Hoc-Tren-Truyen-Hinh-Distance-learning-through-TV-Broadcasting.pdf>. Accessed June 22, 2026.
- [13] World Bank. *E-learning in TVET in Vietnam: Rapid Assessment during COVID-19 and Implications for a National E-learning Platform*. Available at: <https://documents1.worldbank.org/curated/en/0990000005132229736/pdf/P174258045156b0060972302fec6a3015b8.pdf>. Accessed June 22, 2026.
- [14] Honda, D., Halla, K., Thang, S., and Mazari, H. *EdTech in Vietnam: A Rapid Scan*. EdTech Hub, 2024. Available at: <https://docs.edtechhub.org/lib/9RIJIM4K>. Accessed June 22, 2026.